

# Untitled67

December 10, 2025

```
[32]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
from scipy.stats import shapiro
from statsmodels.stats.diagnostic import het_breuschpagan
from statsmodels.stats.stattools import durbin_watson
```

```
[2]: data = pd.read_csv(r"C:\Users\USER\Downloads\Startups.csv")
data.head()
```

```
[2]: R&D Expenditure Administration Expenditure Marketing Expenditure \
```

0	165349.20	136897.80	471784.10
1	162597.70	151377.59	443898.53
2	153441.51	101145.55	407934.54
3	144372.41	118671.85	383199.62
4	142107.34	91391.77	366168.42

	State	Profit
0	Florida	192261.83
1	Florida	191792.06
2	Florida	191050.39
3	Florida	182901.99
4	Florida	166187.94

```
[3]: data.shape
```

```
[3]: (50, 5)
```

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype  

```

```
---  
0    R&D Expenditure      50 non-null      float64  
1    Administration Expenditure 50 non-null      float64  
2    Marketing Expenditure   50 non-null      float64  
3    State                  50 non-null      object  
4    Profit                 50 non-null      float64  
dtypes: float64(4), object(1)  
memory usage: 2.1+ KB
```

```
[5]: data.duplicated().sum()
```

```
[5]: np.int64(0)
```

```
[6]: data.isnull().sum()
```

```
[6]: R&D Expenditure      0  
Administration Expenditure 0  
Marketing Expenditure     0  
State                      0  
Profit                     0  
dtype: int64
```

```
[7]: data.describe()
```

```
[7]:      R&D Expenditure  Administration Expenditure  Marketing Expenditure \\\ncount      50.000000          50.000000          50.000000  
mean      73721.615600        121344.639600        211025.097800  
std       45902.256482        28017.802755        122290.310726  
min       0.000000           51283.140000         0.000000  
25%      39936.370000        103730.875000        129300.132500  
50%      73051.080000        122699.795000        212716.240000  
75%      101602.800000        144842.180000        299469.085000  
max      165349.200000        182645.560000        471784.100000  
  
      Profit  
count      50.000000  
mean     112012.639200  
std      40306.180338  
min      14681.400000  
25%      90138.902500  
50%      107978.190000  
75%      139765.977500  
max      192261.830000
```

```
[8]: y = data['Profit']  
x = data[['Administration Expenditure', 'Marketing Expenditure', 'R&D  
Expenditure']]
```

```
[9]: df = data.select_dtypes(include=['number'])
correlation = df.corr()
print(correlation)
```

	R&D Expenditure	Administration Expenditure	Marketing Expenditure	Profit
R&D Expenditure	1.000000	0.241955	0.724248	0.972900
Administration Expenditure	0.241955	1.000000	-0.032154	0.200717
Marketing Expenditure	0.724248	-0.032154	1.000000	0.747766
Profit	0.972900	0.200717	0.747766	1.000000

	Marketing Expenditure	Profit
R&D Expenditure	0.724248	0.972900
Administration Expenditure	-0.032154	0.200717
Marketing Expenditure	1.000000	0.747766
Profit	0.747766	1.000000

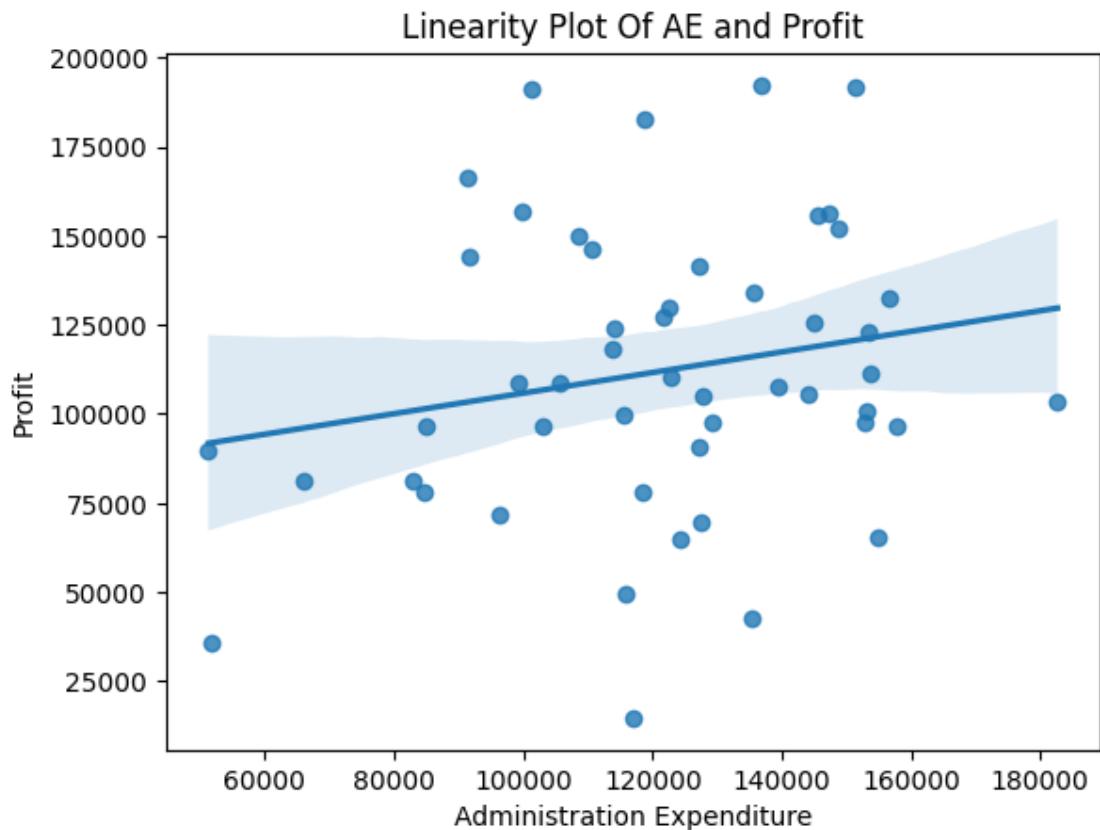
```
[10]: Heatmap = sns.heatmap(correlation, annot= True)
plt.show()
```

<IPython.core.display.Javascript object>



```
[11]: sns.regplot(x='Administration Expenditure', y='Profit', data=data)
plt.title('Linearity Plot Of AE and Profit')
plt.show()
```

<IPython.core.display.Javascript object>



```
[12]: sns.regplot(x='Marketing Expenditure', y='Profit', data=data)
plt.title('Linearity Plot Of ME and Profit')
plt.show()
```

<IPython.core.display.Javascript object>



```
[13]: sns.regplot(x='R&D Expenditure', y='Profit', data=data)
plt.title('Linearity Plot Of R&D and Profit')
plt.show()
```

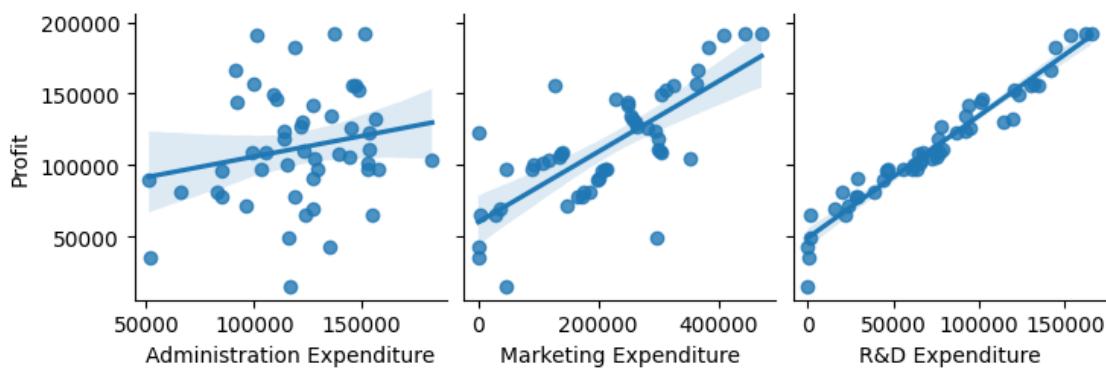
<IPython.core.display.Javascript object>



```
[17]: sns.pairplot(df, x_vars=x.columns, y_vars='Profit', kind='reg')
```

<IPython.core.display.Javascript object>

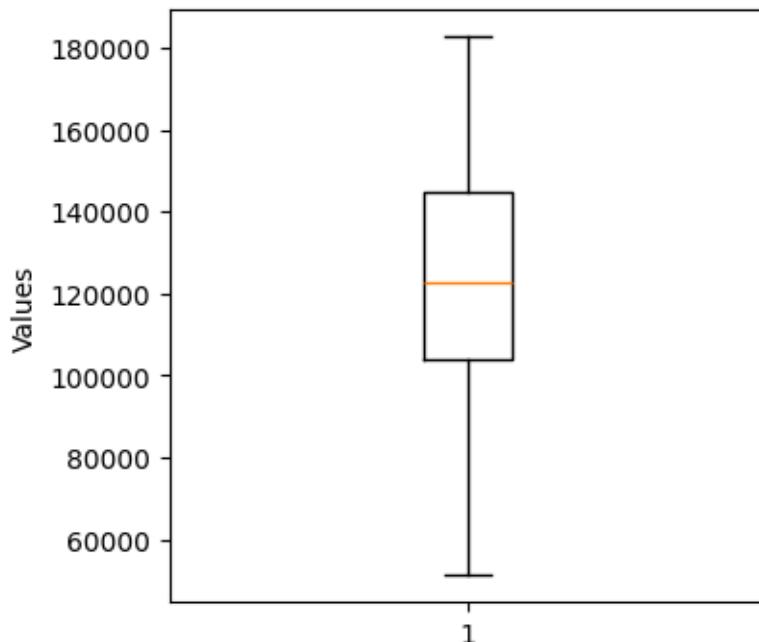
```
[17]: <seaborn.axisgrid.PairGrid at 0x17d1a15d940>
```

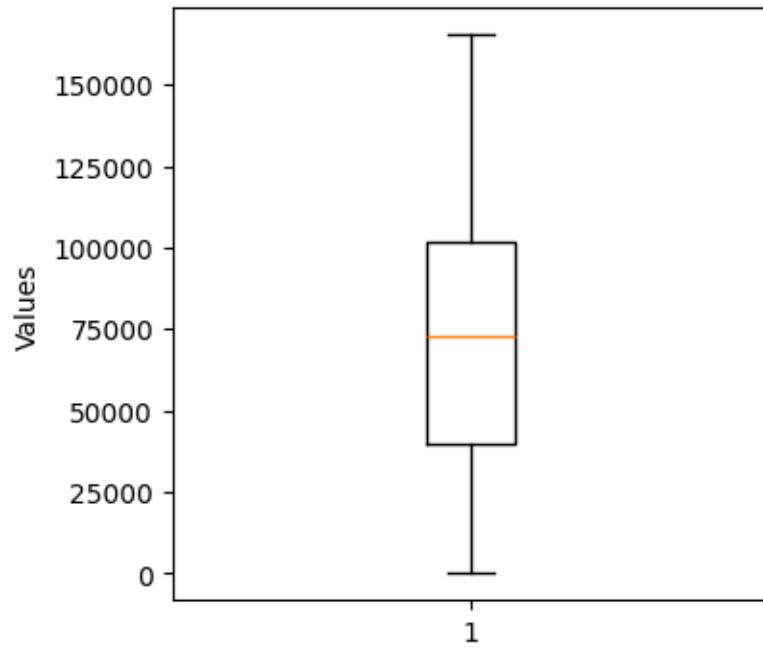
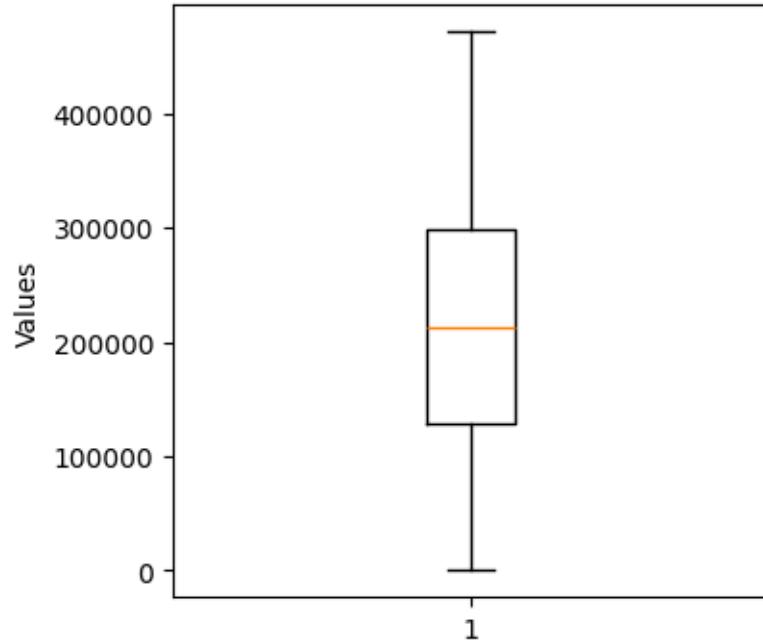


```
[19]: X_const = sm.add_constant(x)
vif = pd.DataFrame({
    'Feature': X_const.columns,
    'VIF': [variance_inflation_factor(X_const.values, i)
             for i in range(X_const.shape[1])]}
)
print(vif)
```

	Feature	VIF
0	const	25.338950
1	Administration Expenditure	1.175091
2	Marketing Expenditure	2.326773
3	R&D Expenditure	2.468903

```
[20]: for col in x.columns:
    plt.figure(figsize=(4, 4))
    plt.boxplot(x[col], vert=True)
    plt.ylabel('Values')
    plt.show()
```





```
[21]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.  
         ↪2,random_state=42)
```

```
[22]: model = LinearRegression()
model.fit(x_train,y_train)
```

```
[22]: LinearRegression()
```

```
[24]: y_pred_train = model.predict(x_train)
y_pred_test = model.predict(x_test)
```

```
[27]: print(r2_score(y_test, y_pred_test))
print(mean_squared_error(y_test, y_pred_test))
print(mean_absolute_error(y_test, y_pred_test))
```

0.900065308303732

80926321.22295165

6979.1522523704025

```
[28]: residuals_train = y_train - y_pred_train
residuals_test = y_test - y_pred_test
```

```
[31]: stat, p = shapiro(residuals_train)
p
```

```
[31]: np.float64(0.013342603561094615)
```

```
[34]: dw = durbin_watson(residuals_train)
dw
```

```
[34]: np.float64(1.7593850382807832)
```

```
[36]: X_train_const = sm.add_constant(x_train)
bp = het_breuschpagan(residuals_train, X_train_const)
bp
```

```
[36]: (np.float64(2.301709808738015),
       np.float64(0.5121934321518875),
       np.float64(0.7326729558482282),
       np.float64(0.5393370968568404))
```

```
[ ]:
```